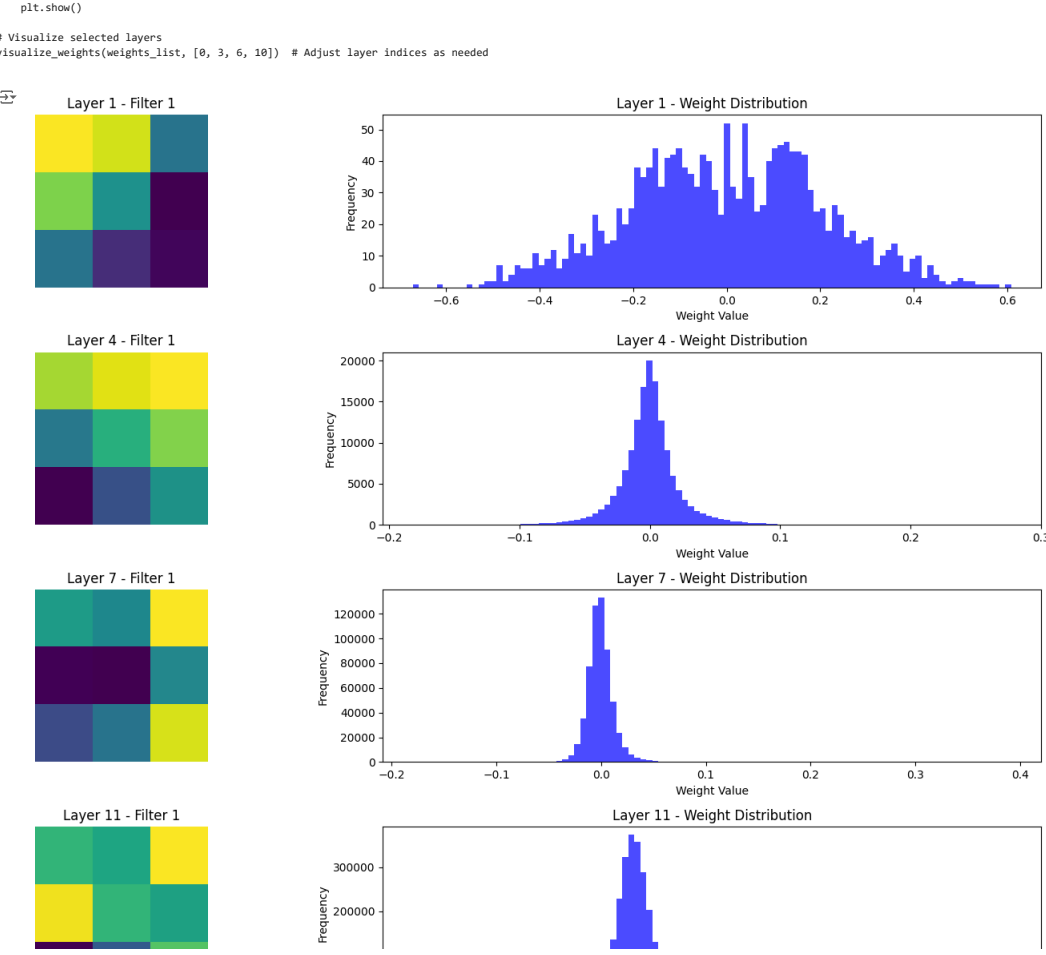



```
layer_weights = w
layer_biases = b

# Layer 1: Input layer
layer_weights = w
layer_biases = b

# Layer 2: Hidden layer
layer_weights = w
layer_biases = b

# Layer 3: Output layer
layer_weights = w
layer_biases = b
```

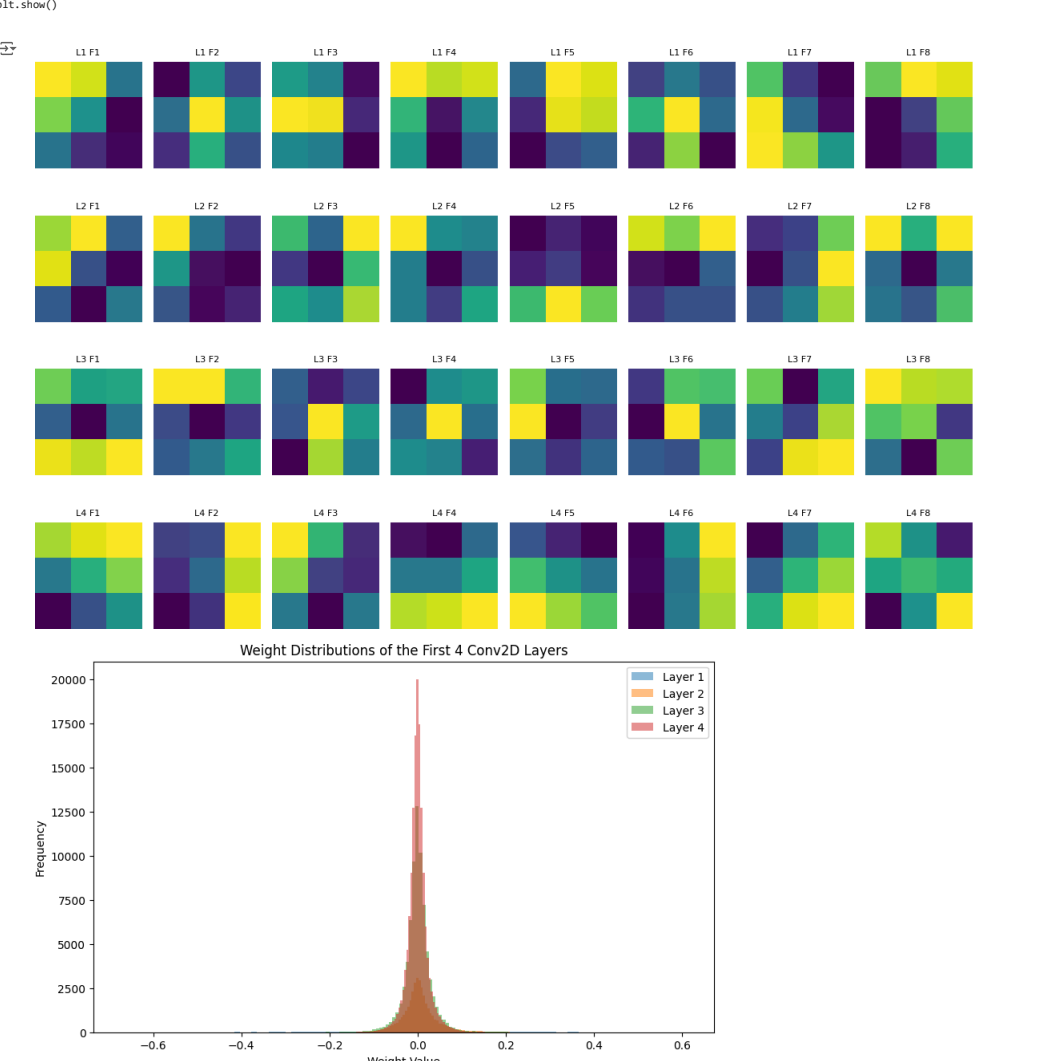


layer_weights = w
layer_biases = b

Layer 1: Input layer
layer_weights = w
layer_biases = b

Layer 2: Hidden layer
layer_weights = w
layer_biases = b

Layer 3: Output layer
layer_weights = w
layer_biases = b




We can also visualize the activation values in the network. We'll do this by applying a forward pass with a color input, and extracting the activation values. Below is an example code for the forward pass.

```
def forward_pass(image):
    # Layer 1: Input layer
    layer_weights = w
    layer_biases = b

    # Layer 2: Hidden layer
    layer_weights = w
    layer_biases = b

    # Layer 3: Output layer
    layer_weights = w
    layer_biases = b
```

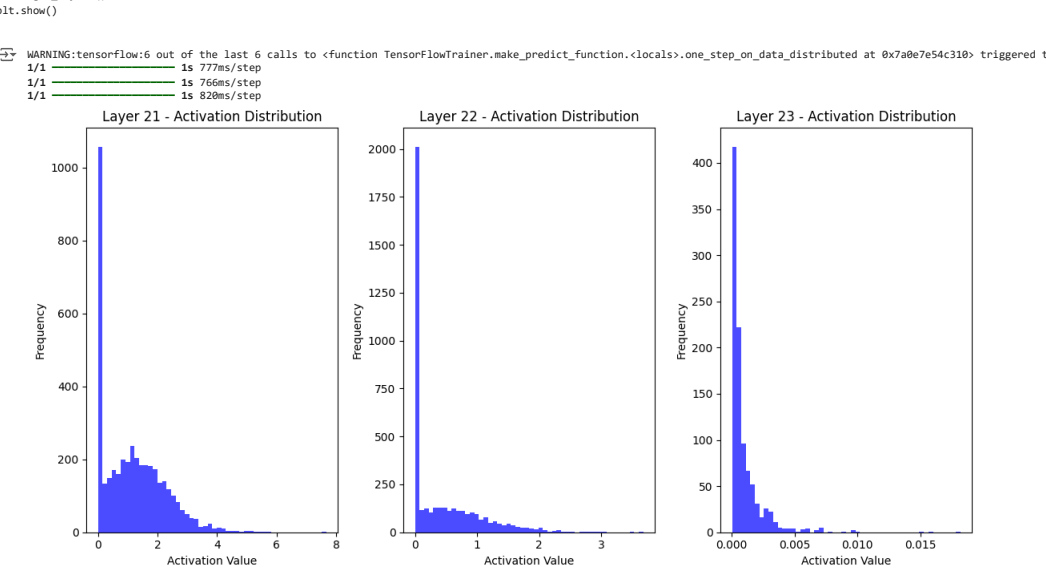


Using the above code for the forward pass, and the layer weights, plot the activation distributions for the first three dense layers.

```
def forward_pass(image):
    # Layer 1: Input layer
    layer_weights = w
    layer_biases = b

    # Layer 2: Hidden layer
    layer_weights = w
    layer_biases = b

    # Layer 3: Output layer
    layer_weights = w
    layer_biases = b
```



What do you notice about the distributions, and how they compare to those of the weight tensors?

Activation distributions are typically more peaked and narrower than weight distributions, indicating that the network is more sensitive to changes in the input data. This is likely due to the non-linear activation functions used in the network, which can cause the output to be more sensitive to changes in the input.